

Probability in Computing

LECTURE 5: MORE APPLICATIONS WITH PROBABILISTIC ANALYSIS, BINS AND BALLS

Agenda

- ◆ Review: Coupon Collector's problem and Packet Sampling
- ◆ Analysis of Quick-Sort
- ◆ Birthday Paradox and applications
- ◆ The Bins and Balls Model


Coupon Collector Problem

- ◆ Problem: Suppose that each box of cereal contains one of n different coupons. Once you obtain one of every type of coupon, you can send in for a prize.
- ◆ Question: How many boxes of cereal must you buy before obtaining at least one of every type of coupon.
- ◆ Let X be the number of boxes bought until at least one of every type of coupon is obtained.
- ◆ $E[X] = nH(n) = n \ln n$

Application: Packet Sampling


- ◆ Sampling packets on a router with probability p
 - The number of packets transmitted after the last sampled packet until and including the next sampled packet is geometrically distributed.
- ◆ From the point of destination host, determining all the routers on the path is like a coupon collector's problem.
- ◆ If there's n routers, then the expected number of packets arrived before destination host knows all of the routers on the path $= n \ln(n)$.

DoS attack

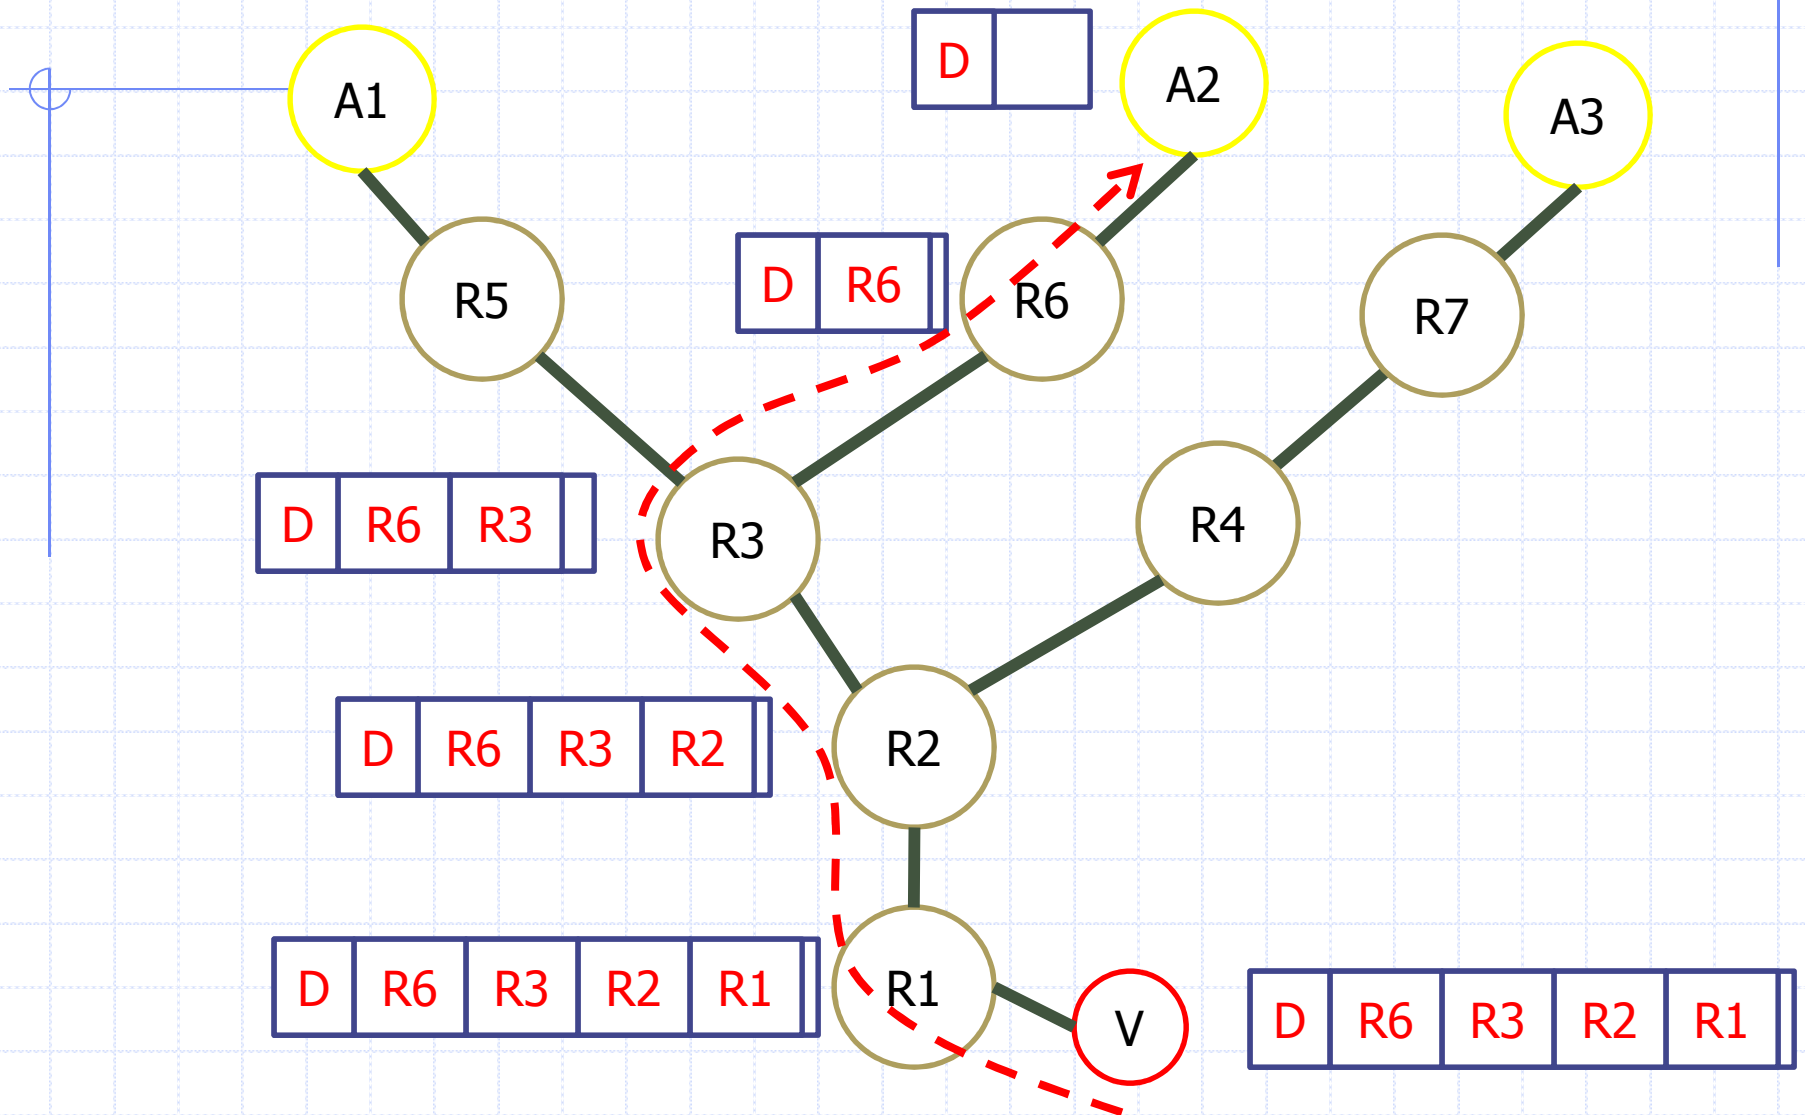


IP traceback

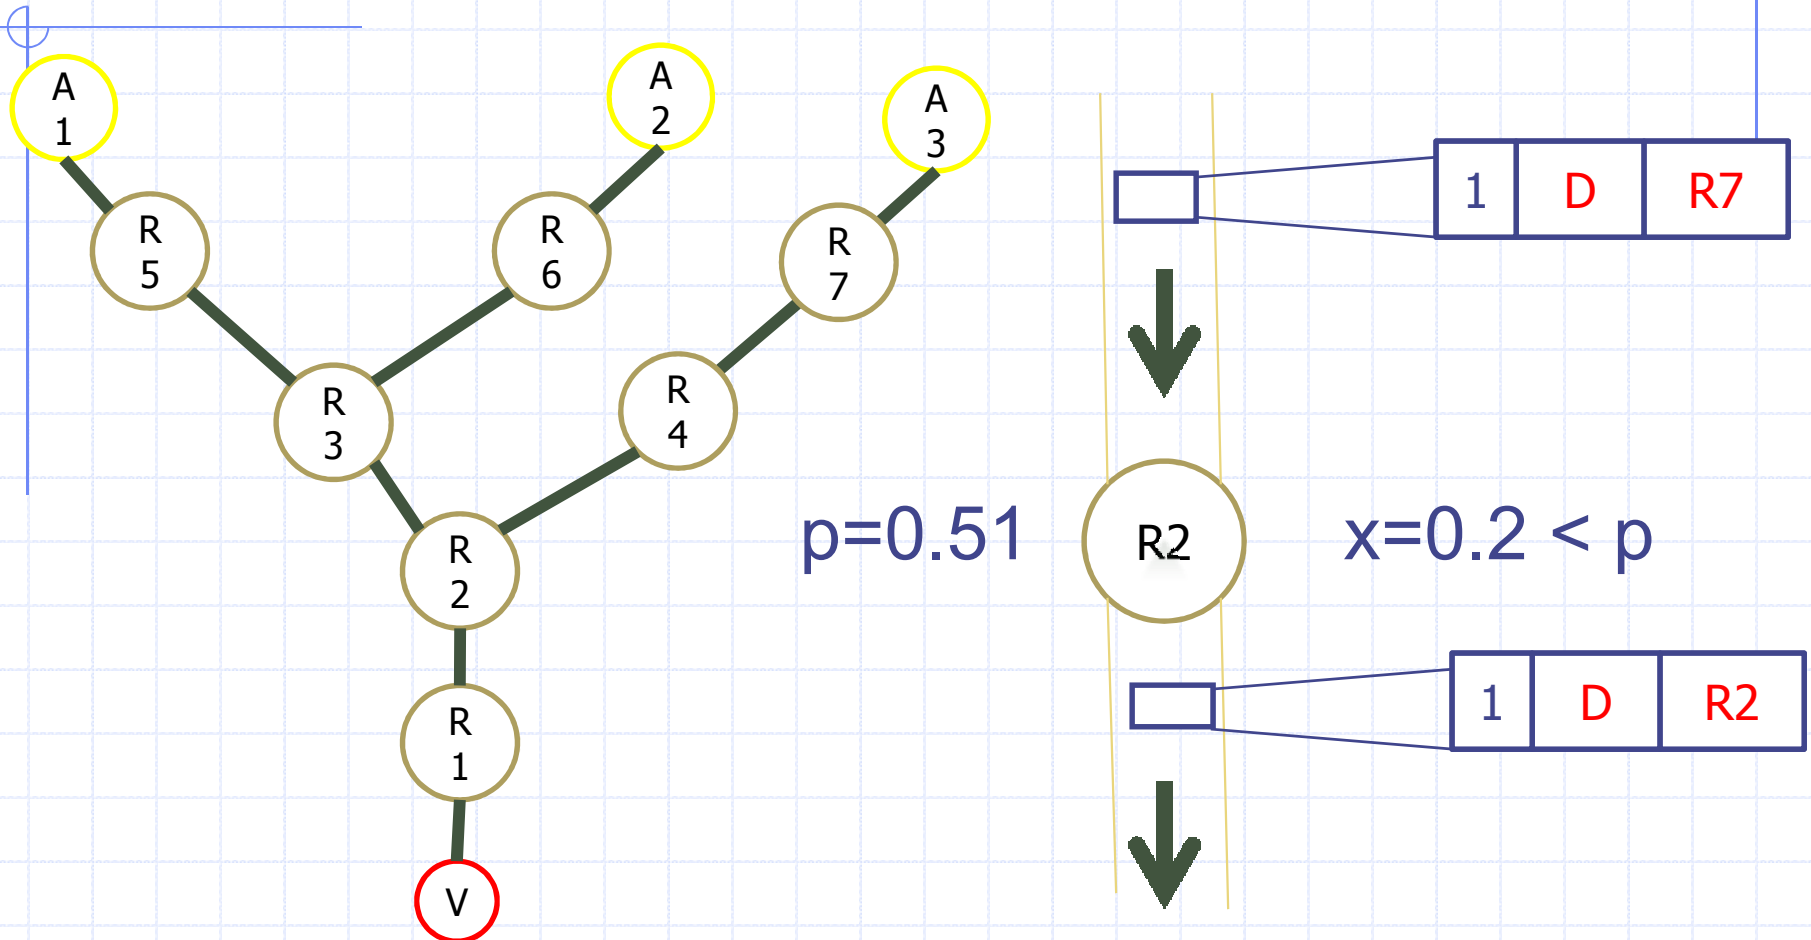
- ◆ Marking and Reconstruction
 - Node append vs. node sampling



Node append



Node Sampling



Expected Run-Time of QuickSort

Quicksort Algorithm:

Input: A list $S = \{x_1, \dots, x_n\}$ of n distinct elements over a totally ordered universe.

Output: The elements of S in sorted order.

1. If S has one or zero elements, return S . Otherwise continue.
2. Choose an element of S as a pivot; call it x .
3. Compare every other element of S to x in order to divide the other elements into two sublists:
 - (a) S_1 has all the elements of S that are less than x ;
 - (b) S_2 has all those that are greater than x .
4. Use Quicksort to sort S_1 and S_2 .
5. Return the list S_1, x, S_2 .

Analysis

- ◆ Worst-case: n^2 .
- ◆ Depends on how we choose the pivot.
- ◆ Good pivot (divide the list in two nearly equal length sub-lists) vs. Bad pivot.
- ◆ In case of good pivot $\rightarrow n \lg(n)$. [by solving recurrence]
- ◆ If we choose pivot point randomly, we will have a randomized version of QuickSort.

Analysis

◆ X_{ij} be a random variable that

- Takes value 1 if y_i and y_j are compared with each other
- 0 if they are not compared.

◆ $E[X] = \sum \sum E[X_{ij}]$

◆ $E[X_{ij}] = 2 / (j - i + 1)$

- Consider when the set $Y_{ij} = \{y_i, y_{i+1}, \dots, y_j\}$ is “touched” by a pivot the first time. If this pivot is either y_i or y_j then the two will be compared, otherwise Never – a (S_1, S_2) split

◆ Using $k = j - i + 1$, we can compute $E[X] = 2n \ln(n)$

Detail analysis

$$\begin{aligned}\mathbf{E}[X] &= \sum_{i=1}^{n-1} \sum_{j=i+1}^n \frac{2}{j-i+1} \\&= \sum_{i=1}^{n-1} \sum_{k=2}^{n-i+1} \frac{2}{k} \\&= \sum_{k=2}^n \sum_{i=1}^{n+1-k} \frac{2}{k} \\&= \sum_{k=2}^n (n+1-k) \frac{2}{k} \\&= \left((n+1) \sum_{k=2}^n \frac{2}{k} \right) - 2(n-1) \\&= (2n+2) \sum_{k=1}^n \frac{1}{k} - 4n.\end{aligned}$$

Birthday “Paradox”

What is the probability that two persons in a room of 30 have the same birthday?

Birthday Paradox

◆ Ways to assign k different birthdays without duplicates:

$$\begin{aligned} N &= 365 * 364 * \dots * (365 - k + 1) \\ &= 365! / (365 - k)! \end{aligned}$$

◆ Ways to assign k different birthdays with possible duplicates:

$$D = 365 * 365 * \dots * 365 = 365^k$$

Birthday “Paradox”

Assuming real birthdays assigned randomly:

N/D = probability there are no duplicates

$1 - N/D$ = probability there is a duplicate

$$= 1 - 365! / ((365 - k)!(365)^k)$$

Generalizing Birthdays

$$P(n, k) = 1 - n! / (n-k)! n^k$$

Given k random selections from n possible values, $P(n, k)$ gives the probability that there is at least 1 duplicate.

Birthday Probabilities

$$P(\text{no two match}) = 1 - P(\text{all are different})$$

$$P(2 \text{ chosen from } N \text{ are different})$$

$$= 1 - 1/N$$

$$P(3 \text{ are all different})$$

$$= (1 - 1/N)(1 - 2/N)$$

$$P(n \text{ trials are all different})$$

$$= (1 - 1/N)(1 - 2/N) \dots (1 - (n - 1)/N)$$

$$\ln(P)$$

$$= \ln(1 - 1/N) + \ln(1 - 2/N) + \dots \ln(1 - (k - 1)/N)$$

Happy Birthday Bob!

$$\ln(P) = \ln(1 - 1/N) + \dots + \ln(1 - (k-1)/N)$$

For $0 < x < 1$: $\ln(1 - x) \leq -x$

$$\ln(P) \leq -(1/N + 2/N + \dots + (k-1)/N)$$

Gauss says:

$$1 + 2 + 3 + 4 + \dots + (k-1) + k = \frac{1}{2} k(k+1)$$

So,

$$\ln(P) \leq -\frac{1}{2} (k-1) k / N$$

$$P \leq e^{-\frac{1}{2} (k-1) k / N}$$

$$\text{Probability of match} \geq 1 - e^{-\frac{1}{2} (k-1) k / N}$$

Applying Birthdays

$$P(n, k) > 1 - e^{-k*(k-1)/2n}$$

⑩ For $n = 365, k = 20$:

$$P(365, 20) > 1 - e^{-20*(19)/2*365}$$

$$P(365, 20) > .4058$$

⑩ For $n = 2^{64}, k = 2^{32}$: $P(2^{64}, 2^{32}) > .39$

⑩ For $n = 2^{64}, k = 2^{33}$: $P(2^{64}, 2^{33}) > .86$

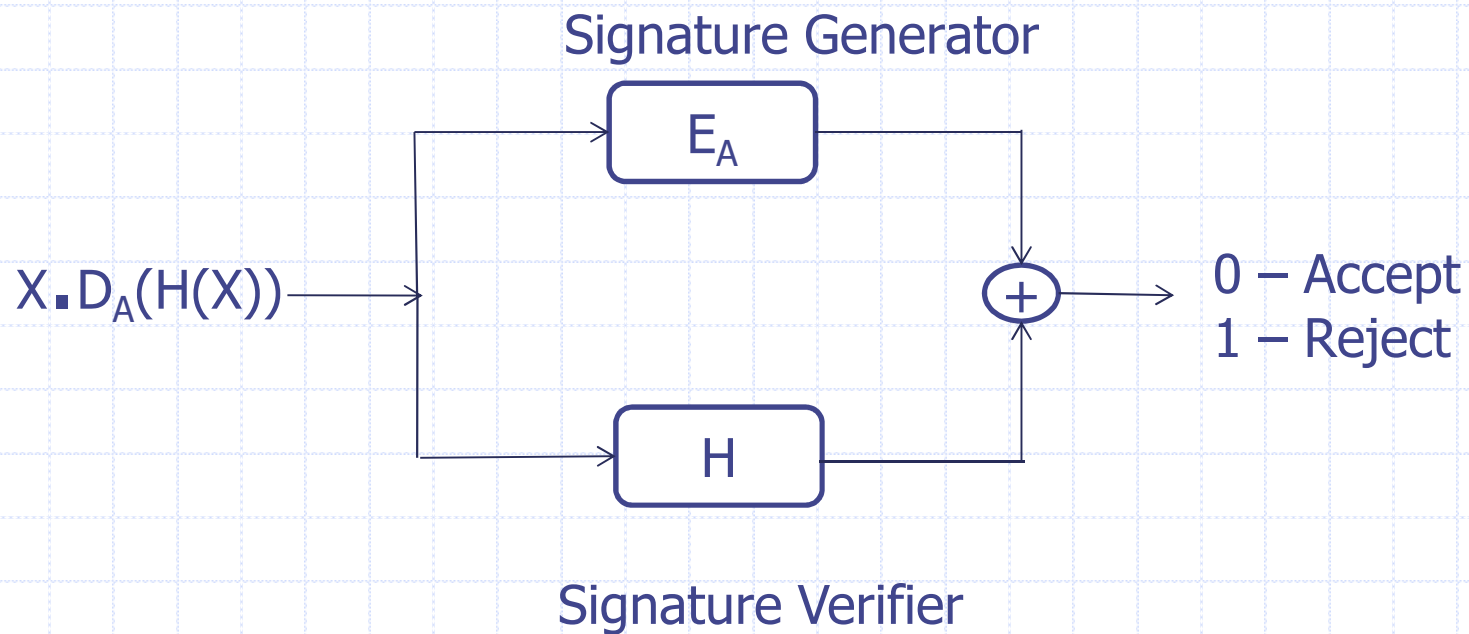
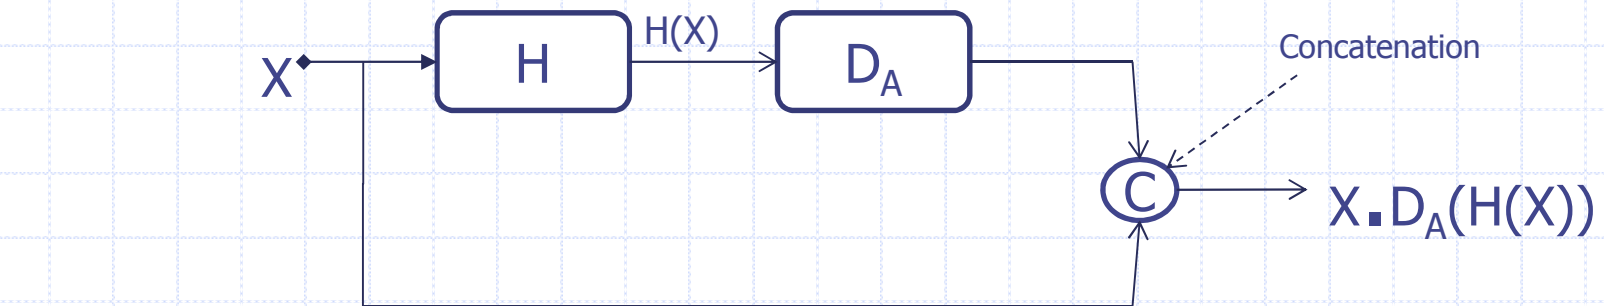
⑩ For $n = 2^{64}, k = 2^{34}$: $P(2^{64}, 2^{34}) > .99996$

⑩ Application: Digital Signatures

Digital Signature Scheme: Using Hash Functions

- ◆ A hash function H maps a message of variable length n bits to a fingerprint of fixed length m bits, with $m < n$.
 - This hash value is also called a digest (of the original message).
 - Since $n > m$, there exist many X which are map to the same digest → collision.

DS schemes with hash functions



Main properties

Given a hash function $H: X \rightarrow Y$

- ◆ Long message \rightarrow short, fixed-length hash

- ◆ One-way property: given $y \in Y$

it is computationally infeasible to find a value $x \in X$
s.t. $H(x) = y$

- ◆ Collision resistance (collision-free)

it is computationally infeasible to find any two
distinct values $x', x \in X$ s.t. $H(x') = H(x)$

- This property prevent against signature forgery

Collisions

- ◆ Avoiding collisions is theoretically impossible
 - Dirichlet principle: $n+1$ rabbits into n cages \rightarrow at least 2 rabbits go to the same cage
 - This suggest exhaustive search: try $|Y|+1$ messages then must find a collision ($H:X \rightarrow Y$)
- ◆ In practice
 - Choose $|Y|$ large enough so exhaustive search is computational infeasible.
 - ◆ $|Y|$ not too large or long signature and slow process
 - However, collision-freeness is still hard

Birthday attack

◆ Can hash values be of 64 bits?

- Look good, initially, since a space of size 2^{64} is too large to do exhaustive search or compute that many hash values
- However a birthday attack can easily break a DS with a 64-bit hash function
 - ◆ In fact, the attacker only need to create a bunch of 2^{32} messages and then launch the attack with reasonably high probability for success.

How is the attack

- ◆ Goal: given H , find x, x' such that $H(x)=H(x')$
- ◆ Algorithm:
 - pick a random set S of q values in X
 - for each $x \in S$, computes $h_x = H(x)$
 - if $h_x = h_{x'}$ for some $x' \neq x$ then collision found: (x, x') , else fail
- ◆ The average success probability is
$$\varepsilon = 1 - \exp(-q(q-1)/2|Y|)$$
 - Suppose Y has size 2^m , choose $q \approx 2^{m/2}$ then ε is almost 0.5!

Balls into Bins

- ◆ We have m balls that are thrown into n bins, with the location of each ball chosen independently and uniformly at random from n possibilities.
- ◆ What does the distribution of the balls into the bins look like
 - “Birthday paradox” question: is there a bin with at least 2 balls
 - How many of the bins are empty?
 - How many balls are in the fullest bin?

Answers to these questions give solutions to many problems in the design and analysis of algorithms

The maximum load

- ◆ When n balls are thrown independently and uniformly at random into n bins, the probability that the maximum load is more than $3 \ln n / \ln \ln n$ is at most $1/n$ for n sufficiently large.

- By Union bound, $\Pr [\text{bin 1 receives } \geq M \text{ balls}] \leq \binom{n}{M} \left(\frac{1}{n}\right)^M$.
- Note that:

$$\binom{n}{M} \left(\frac{1}{n}\right)^M \leq \frac{1}{M!} \leq \left(\frac{e}{M}\right)^M.$$



- Now, using Union bound again, $\Pr [\text{any bin receives } \geq M \text{ balls}]$ is at most

$$n \left(\frac{e}{M}\right)^M \leq n \left(\frac{e \ln \ln n}{3 \ln n}\right)^{3 \ln n / \ln \ln n}$$

which is $\leq 1/n$

Application: Bucket Sort

- ◆ A sorting algorithm that breaks the $\Omega(n \log n)$ lower bound under certain input assumption
- ◆ Bucket sort works as follows:
 - Set up an array of initially empty "buckets."
 - Scatter: Go over the original array, putting each object in its bucket.
 - Sort each non-empty bucket.
 - Gather: Visit the buckets in order and put all elements back into the original array.



- ◆ A set of $n = 2^m$ integers, randomly chosen from $[0, 2^k)$, $k \geq m$, can be sorted in expected time $O(n)$

- Why: will analyze later!

The Poisson Distribution

◆ Consider m balls, n bins

- $\Pr[\text{a given bin is empty}] = \left(1 - \frac{1}{n}\right)^m \approx e^{-m/n}$;
- Let X_j is a indicator r.v. that os 1 if bin j empty, 0 otherwise
- Let X be a r.v. that represents # empty bins

$$\mathbf{E}[X] = \mathbf{E}\left[\sum_{i=1}^n X_i\right] = \sum_{i=1}^n \mathbf{E}[X_i] = n\left(1 - \frac{1}{n}\right)^m \approx ne^{-m/n}$$

- Generalizing this argument, $\Pr[\text{a given bin has } r \text{ balls}] =$

$$\binom{m}{r} \left(\frac{1}{n}\right)^r \left(1 - \frac{1}{n}\right)^{m-r} = \frac{1}{r!} \frac{m(m-1) \cdots (m-r+1)}{n^r} \left(1 - \frac{1}{n}\right)^{m-r}$$

- Approximately, $p_r \approx \frac{e^{-m/n} (m/n)^r}{r!}$

- So: **Definition 5.1:** A discrete Poisson random variable X with parameter μ is given by the following probability distribution on $j = 0, 1, 2, \dots$:

$$\Pr(X = j) = \frac{e^{-\mu} \mu^j}{j!}.$$

Limit of the Binomial Distribution

We have shown that, when throwing m balls randomly into b bins, the probability p_r that a bin has r balls is approximately the Poisson distribution with mean m/b . In general, the Poisson distribution is the limit distribution of the binomial distribution with parameters n and p , when n is large and p is small. More precisely, we have the following limit result.

Theorem 5.5: *Let X_n be a binomial random variable with parameters n and p , where p is a function of n and $\lim_{n \rightarrow \infty} np = \lambda$ is a constant that is independent of n . Then, for any fixed k ,*

$$\lim_{n \rightarrow \infty} \Pr(X_n = k) = \frac{e^{-\lambda} \lambda^k}{k!}.$$

This theorem directly applies to the balls-and-bins scenario. Consider the situation where there are m balls and b bins, where m is a function of b and $\lim_{n \rightarrow \infty} m/b = \lambda$. Let X_n be the number of balls in a specific bin. Then X_n is a binomial random variable with parameters m and $1/b$. Theorem 5.5 thus applies and says that

$$\lim_{n \rightarrow \infty} \Pr(X_n = r) = \frac{e^{-m/n} (m/n)^r}{r!},$$